

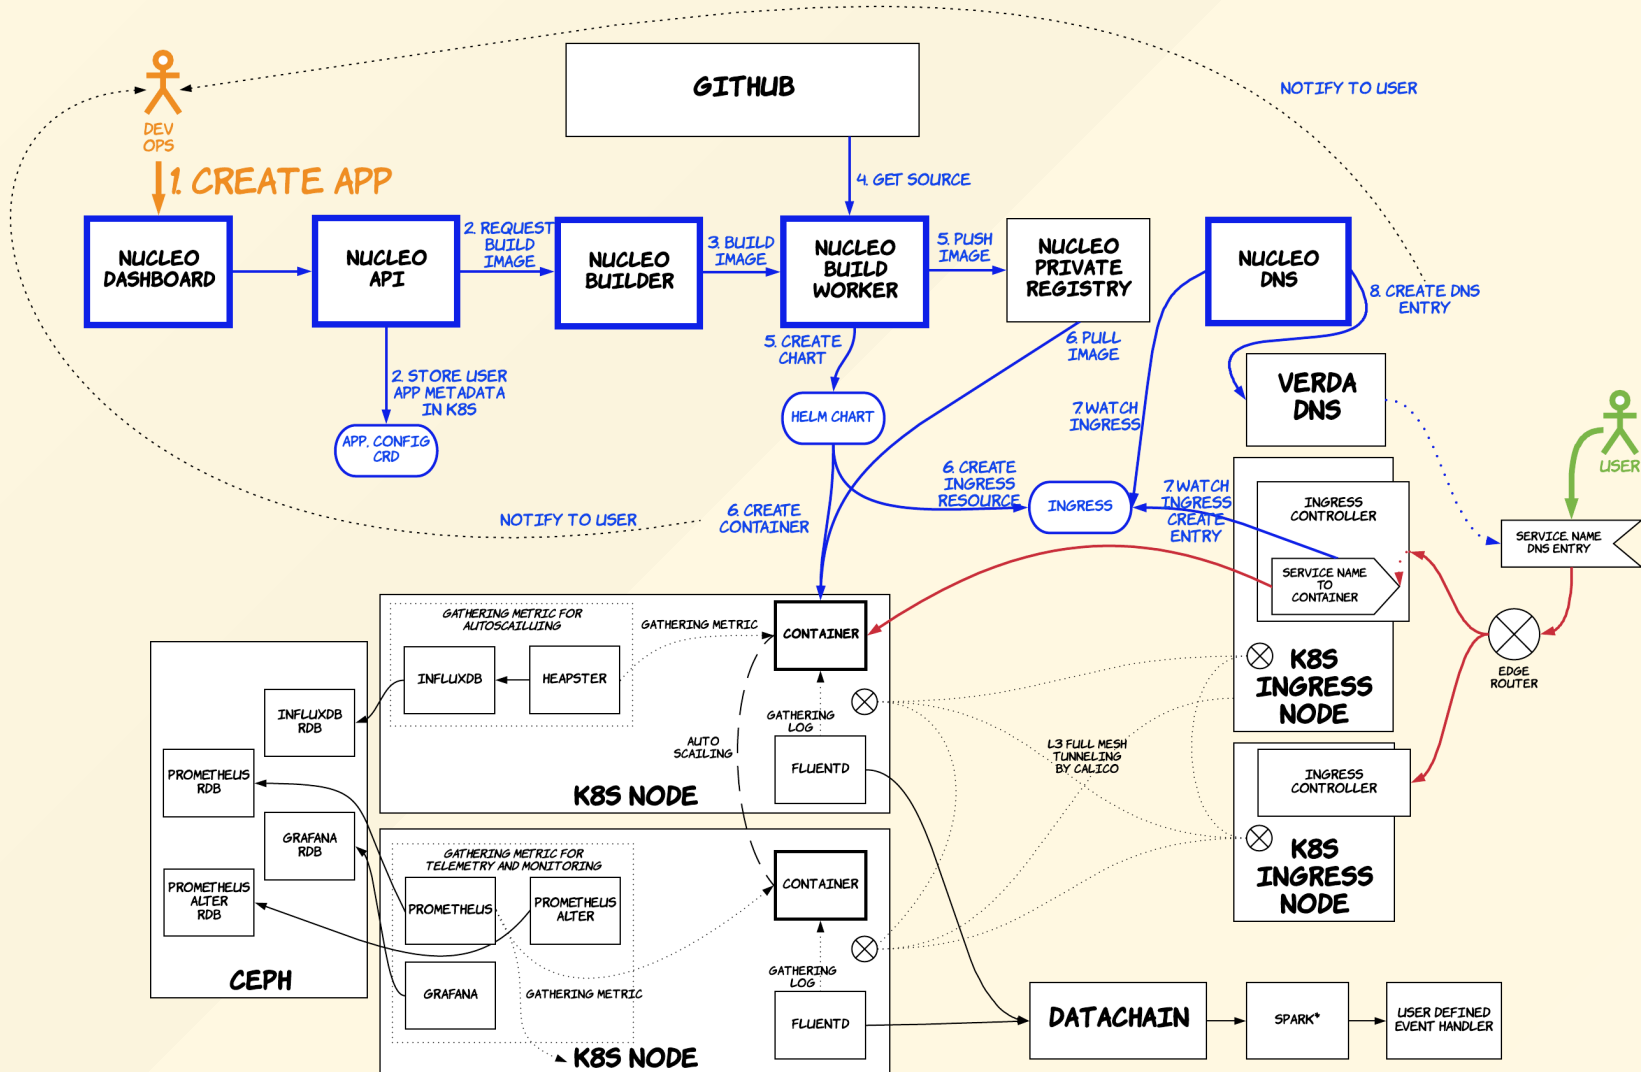
How to make cloud native platform by kubernetes

Eohyung Lee

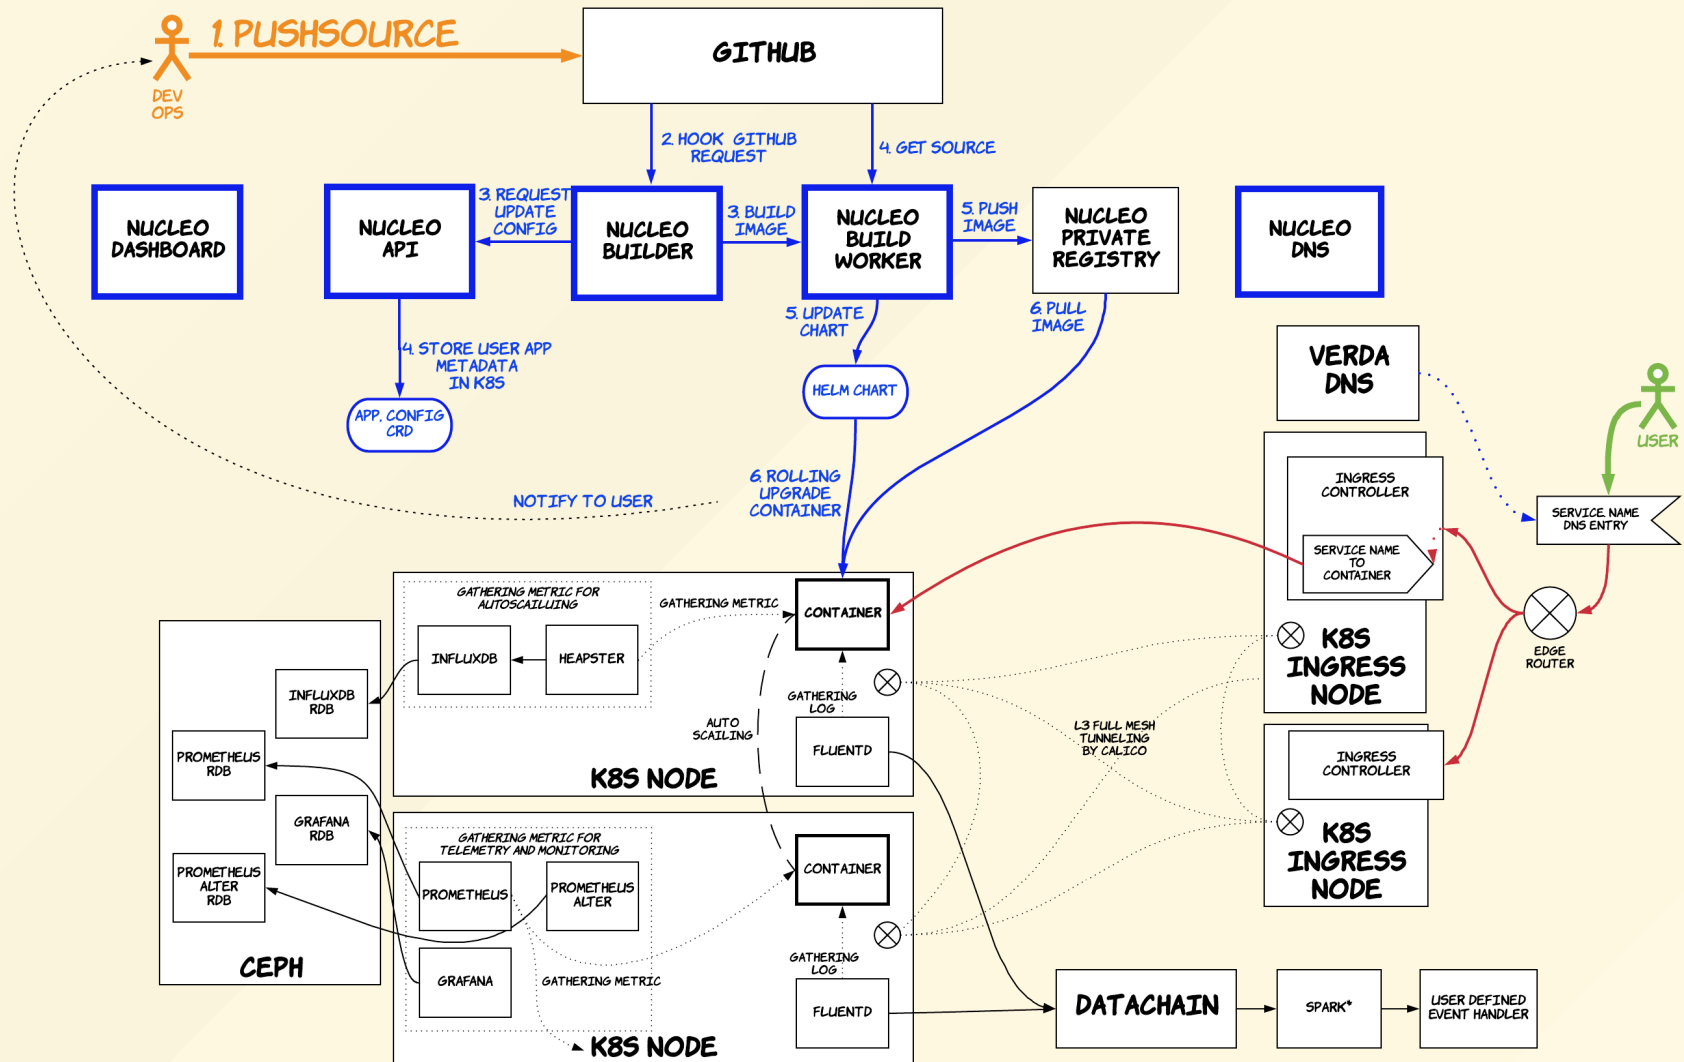
al.lee@linecorp.com

liquidnuker@gmail.com

Nucleo architecture (create app)



Nucleo architecture (update app)



cloud native platform

nucleo

Kubernetes

<https://kubernetes.io/docs/concepts/>

Kubernetes has a number of features. It can be thought of as:

- a container platform
- a microservices platform
- a portable cloud platform and a lot more.

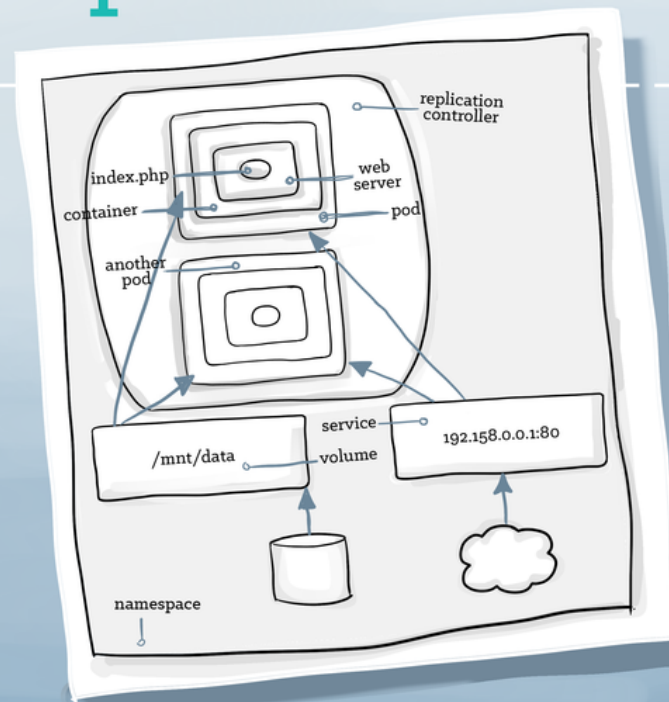
Kubernetes a.k.a. K8S

- $k + \text{len}(\text{ubernetes}) + s = k + 8 + s = k8s$
- provides a `container-centric management environment`.
- orchestrates `computing, networking, and storage infrastructure` on behalf of user workloads.

Traditional explanation of k8s

Namespaces

- Group + segment pods, rcs, volumes, & secrets from each other




The Illustrated Children's Guide to Kubernetes
<https://youtu.be/4ht22ReBjno>

So, we focused to talk about

- how to use k8s as a `cloud native backend platfom`.
 - how to use k8s `by api`
 - How to manage `user app`
 - How to manage `user app metadata`
 - How to support `accessing to public`
 - How to support `continuous delivery`

What constitutes a k8s

- k8s **object** (resource object, **resource**)
 - are **persistent entities** in the Kubernetes system.
 - Kubernetes uses these entities to represent **the state of your cluster.**



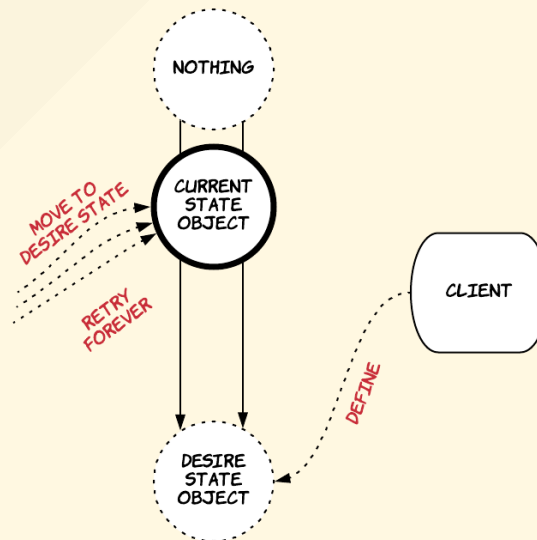
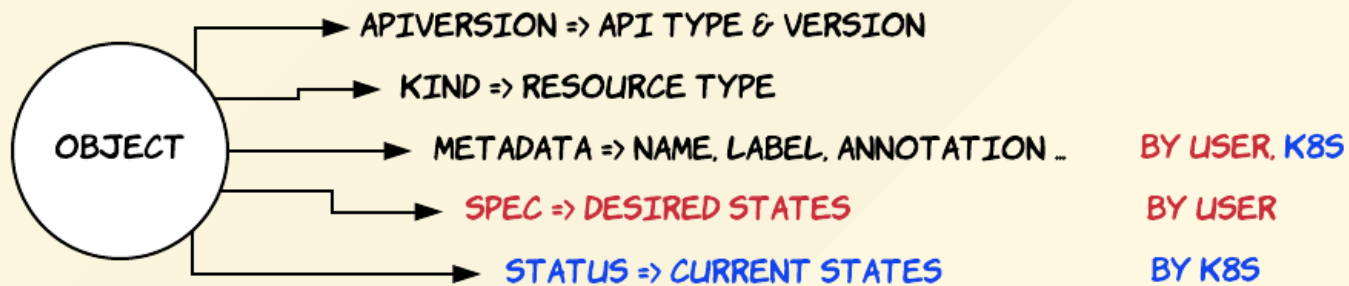
K8S object types

- Workloads
 - container, pods, deployment, daemonset, ...
- Discovery & LB resources
 - endpoint, ingress, service, ...
- Config & Storage resources
 - configmap, secret, persistentvolumeclaim, ...

K8S object types (cont.)

- Cluster resources
 - node, namespace, role, serviceaccount, ...
- Metadata resources
 - event, horizontalpodautoscaler, customresourcedefinition, ...

K8S object components



K8S object spec

- before apply

```
1  apiVersion: autoscaling/v1
2  kind: HorizontalPodAutoscaler
3  metadata:
4    labels:
5      app: api
6      release: nucleo-api
7    name: api
8    namespace: kube-system
9  spec:
10    maxReplicas: 10
11    minReplicas: 3
12    scaleTargetRef:
13      kind: Deployment
14      name: api
15    targetCPUUtilizationPercentage: 20
16
```

K8S object spec & status

- after apply

```
1  apiVersion: autoscaling/v1
2  kind: HorizontalPodAutoscaler
3  metadata:
4    annotations:
5      ...
6    creationTimestamp: 2018-03-12T09:01:48Z
7    labels:
8      app: api
9      release: nucleo-api
10   name: api
11   namespace: kube-system
12   resourceVersion: "179540"
13   selfLink: /apis/autoscaling/v1/namespaces/kube-system/horizontalpodautoscalers/api
14   uid: fef84c9c-25d3-11e8-b71c-fa163ee0bacc
15  spec:
16    maxReplicas: 10
17    minReplicas: 3
18    scaleTargetRef:
19      kind: Deployment
20      name: api
21    targetCPUUtilizationPercentage: 20
22  status:
23    currentCPUUtilizationPercentage: 9
24    currentReplicas: 3
25    desiredReplicas: 3
26
```

Transaction vs spec & status method

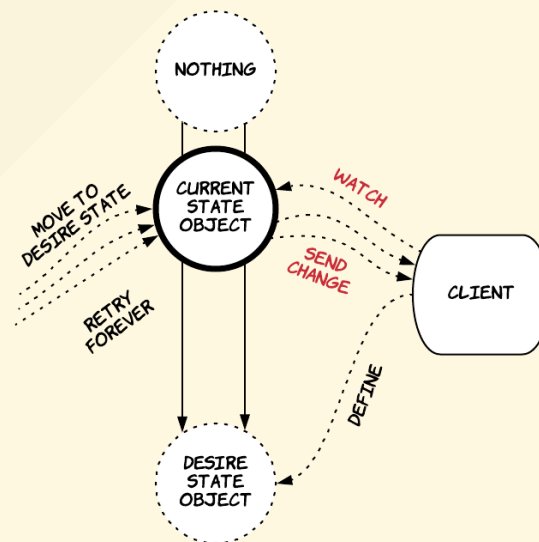
- for supporting micro service architecture
 - infinitely retrievable
 - self healing
 - ...

K8S object operations

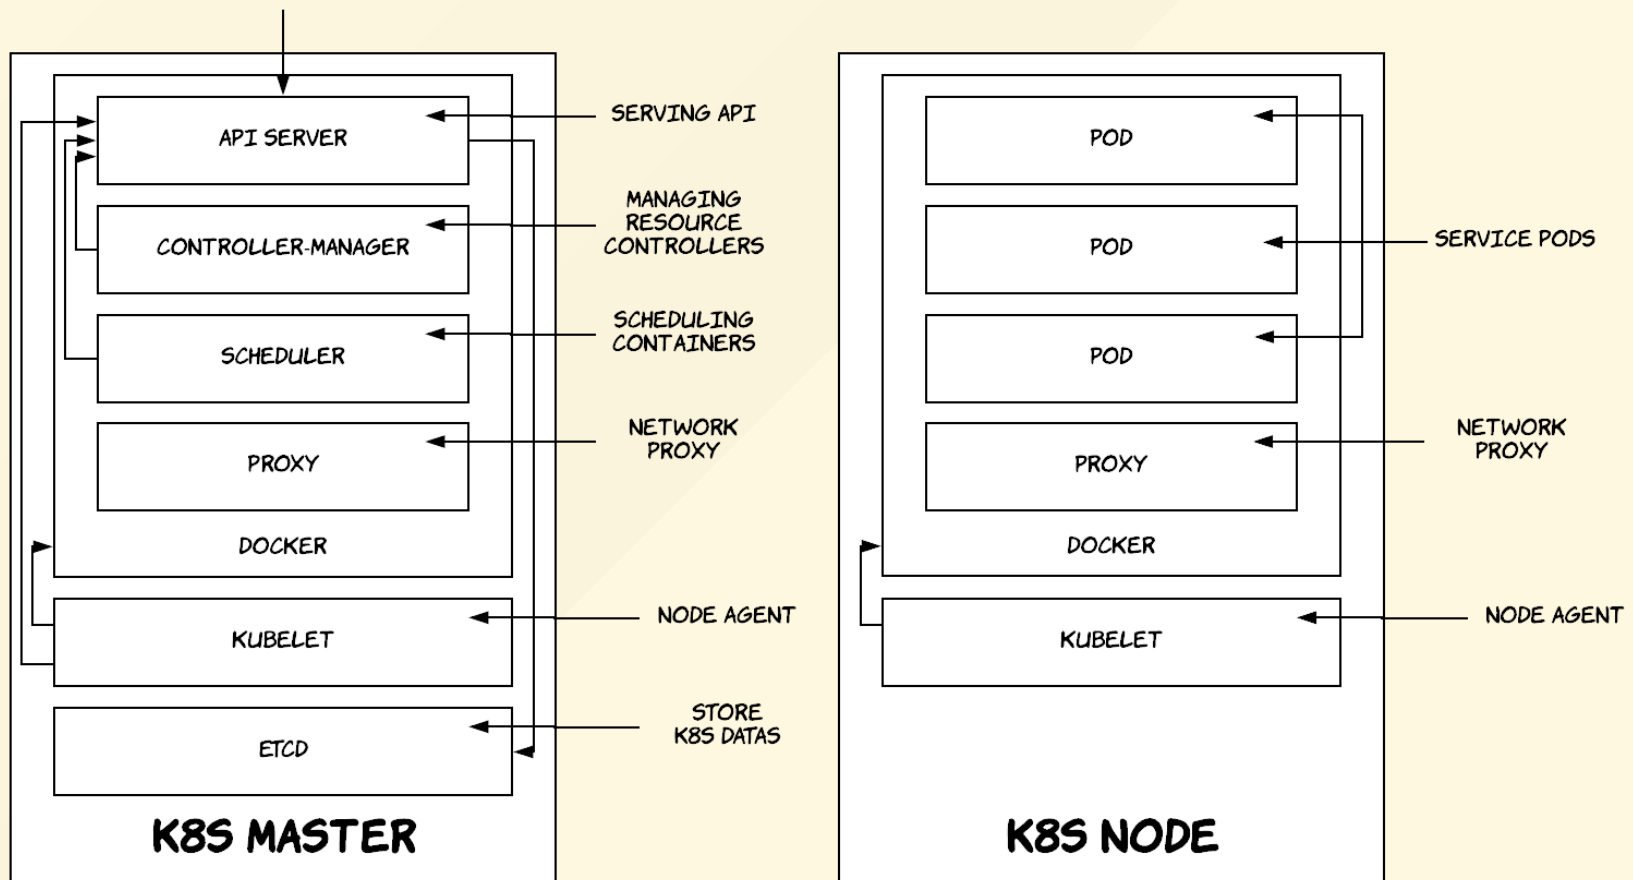
- Create
- Update
 - Replace, Patch
- Read
 - Get, List, Watch
- Delete

K8S watch api

- Watch will **stream results** for an object(s) as it is updated.
- Similar to a callback, watch is used to respond to **resource changes**.




K8S architecture

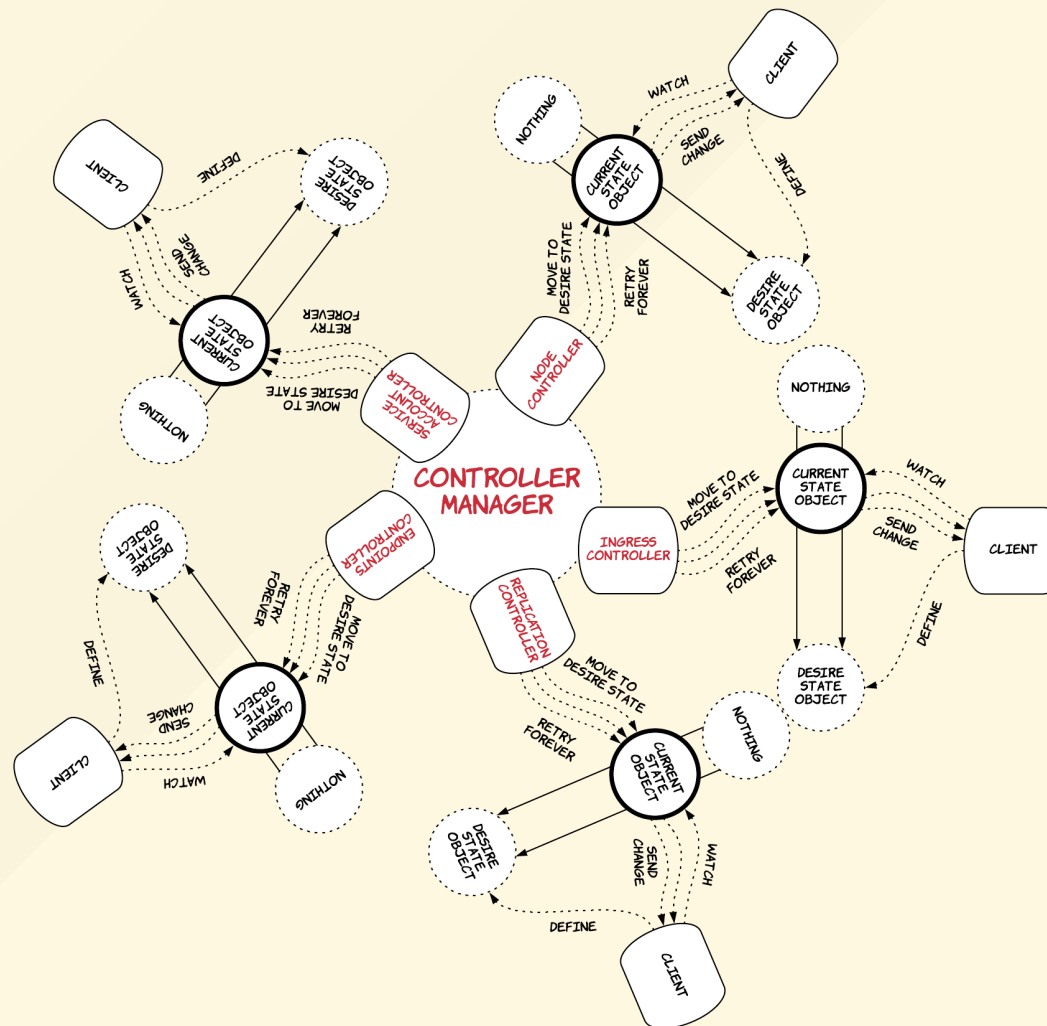


K8S controller

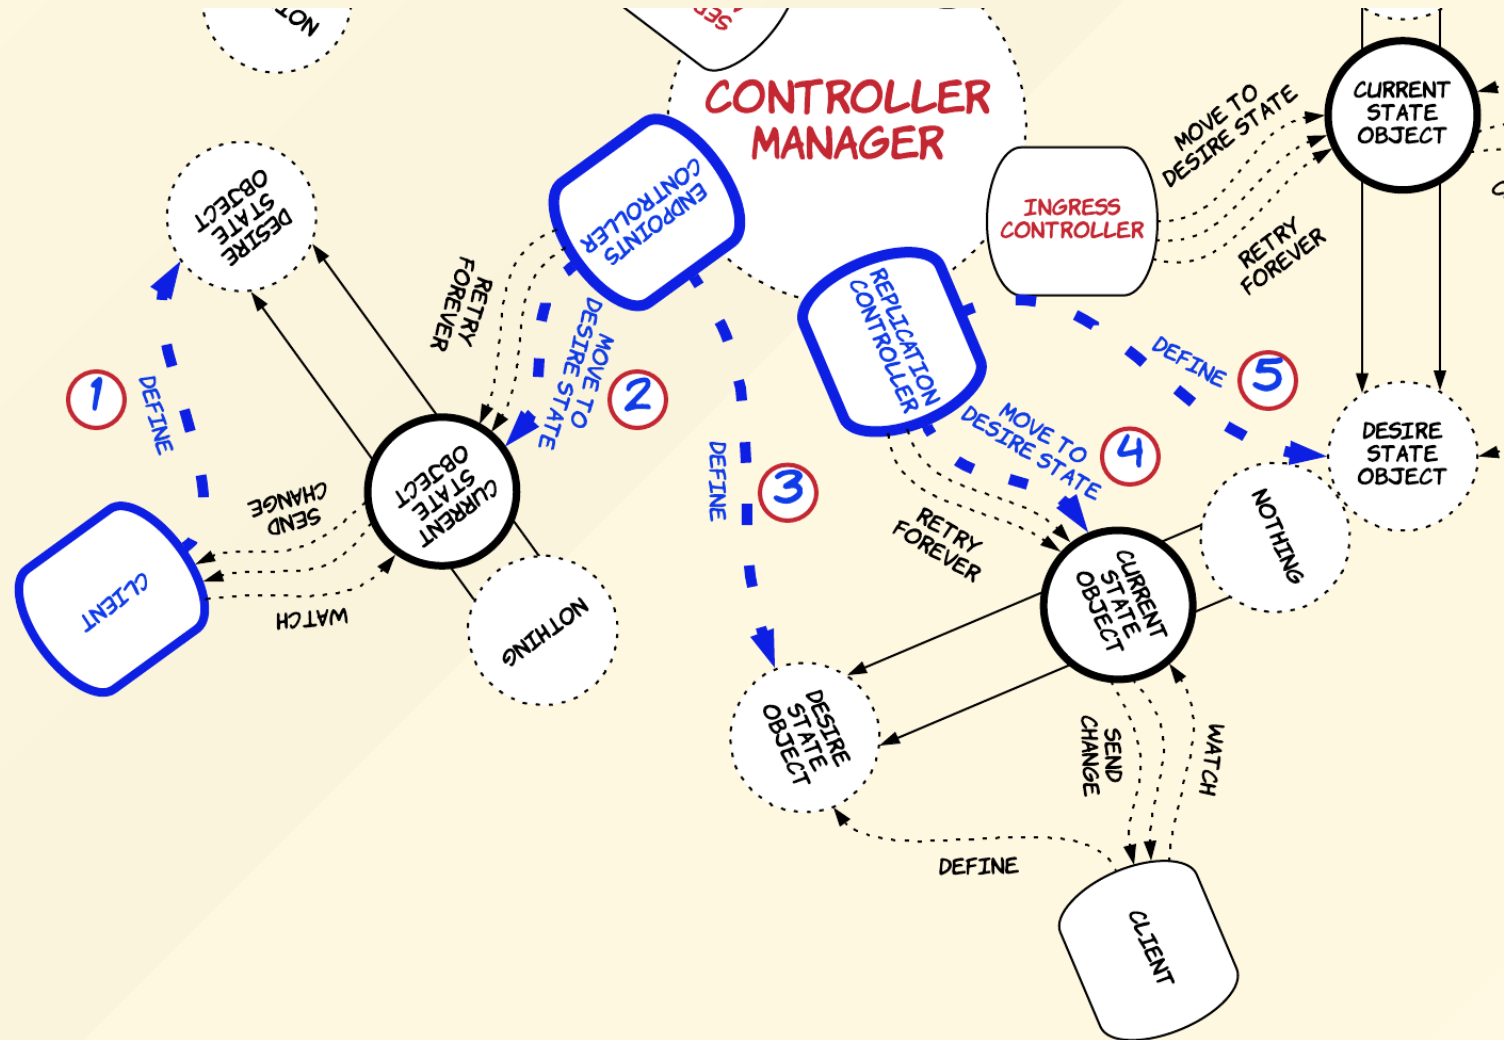
- a controller is a control loop that watches the shared state of the cluster through the apiserver and makes changes attempting to move the **current state** towards the **desired state**



K8S controller manager



For example, chaining events

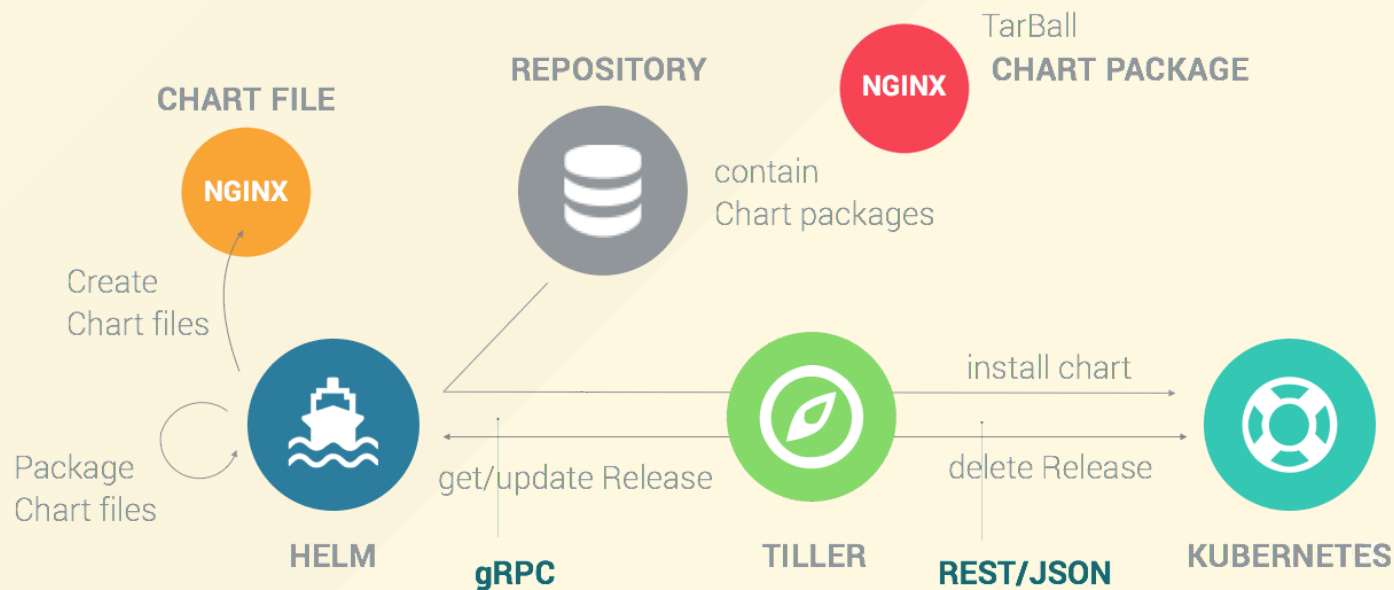


So, k8s API

- remember `spec & status`
- remember `watch`

How to manage user app

- helm: The package manager for Kubernetes



- <https://qiita.com/Ladicle/items/63cad824e27aa8aac7e1>


Nucleo chart

- chart: Helm uses a packaging format called charts

```
charts/nucleo/
├── .helmignore # helm ignore file
├── Chart.yaml # chart spec
├── README.md
├── templates
│   ├── NOTES.txt
│   ├── _helpers.tpl # help function, define values
│   ├── config.yaml # config for user apps
│   ├── deployment.yaml # container deploy method, config
│   ├── hpa.yaml # autoscaling features
│   ├── ingress.yaml # access to outside L7(like nginx)
│   ├── secret.yaml # secret config for user apps
│   └── service.yaml # service discovery, endpoint, L4
└── values.yaml # chart configuration value
```


How to manage user app metadata

- e.g. application information, release histories, build histories ...
- Custom Resource Definition(CRD)
 - store metadata in k8s like other objects



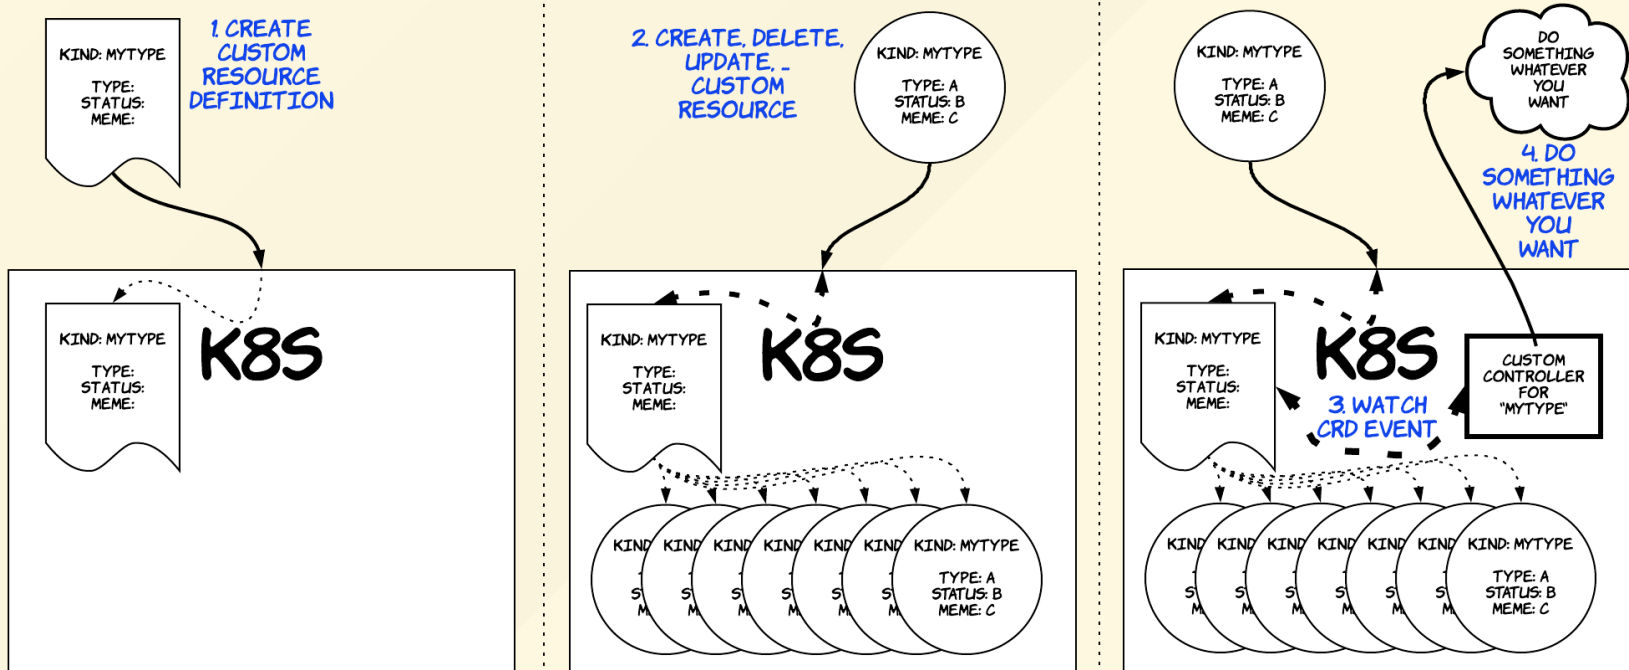
Custom resources

- make your own resource dynamically
- and use k8s database, API, authentication like the other `resource objects`
- support CRUD functions
 - e.g. `Create`, `Update`, `Delete`, `Read`

Custom resources definition

```
apiVersion: nucleo.linecorp.com/v1
kind: DeployConfig # this resource is made by nucleo
metadata:
  creationTimestamp: 2018-01-31T06:03:13Z
  labels:
    nucleo.app: fernet-decryptor
  name: 0162733d-c8e9-44f5-9b3b-f8d587602b17
  resourceVersion: "37313981"
  selfLink: /apis/nucleo.linecorp.com/v1/link
  uid: 6d14f769-064c-11e8-9b84-fa163e9b48b8
spec:
  branch: ^.
  display_name: default
status:
  build_status: success
  status: running
```

Custom resource & controller



How to support accessing to public

- Use ingress for exposing user app to the world
 - like a `aws application load balancer`
 - ingress resources has no prebuilt controller
 - need to use `NGINX Ingress Controller` install by `helm`
 - can support multiple & various ingress controller in one k8s cluster

Ingress with service

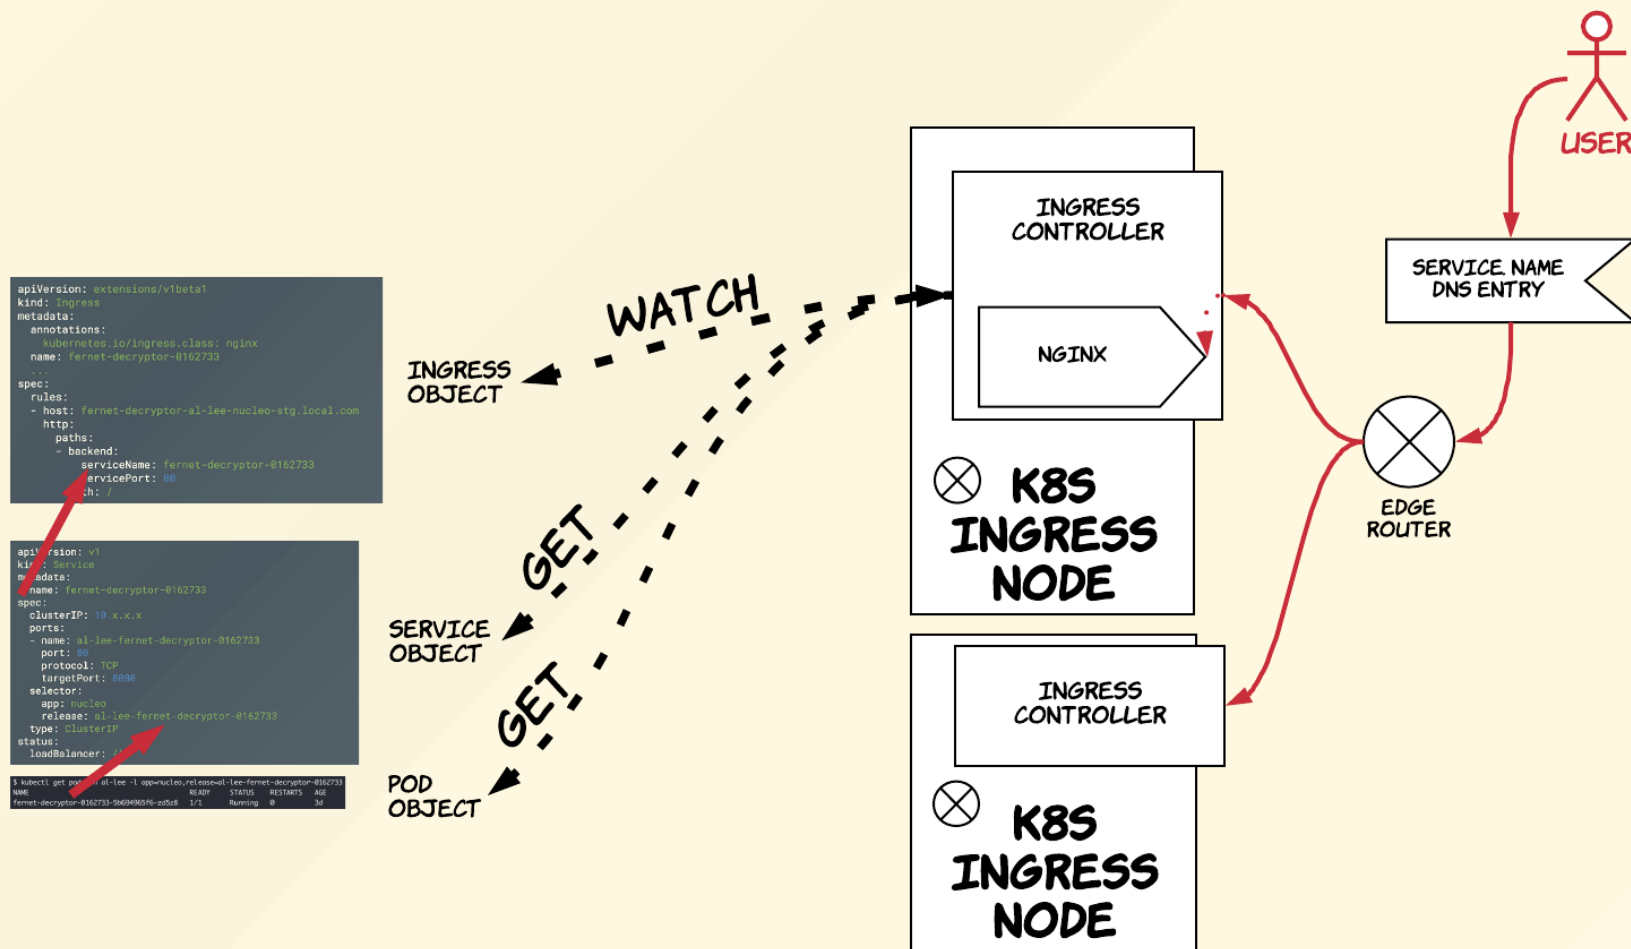
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    kubernetes.io/ingress.class: nginx
  name: fernet-decryptor-0162733
  ...
spec:
  rules:
  - host: fernet-decryptor-al-lee-nucleo-stg.local.com
    http:
      paths:
      - backend:
          serviceName: fernet-decryptor-0162733
          servicePort: 80
        path: /
```

```
apiVersion: v1
kind: Service
metadata:
  name: fernet-decryptor-0162733
spec:
  clusterIP: 10.x.x.x
  ports:
  - name: al-lee-fernet-decryptor-0162733
    port: 80
    protocol: TCP
    targetPort: 8080
  selector:
    app: nucleo
    release: al-lee-fernet-decryptor-0162733
  type: ClusterIP
  status:
    loadBalancer: {}
```

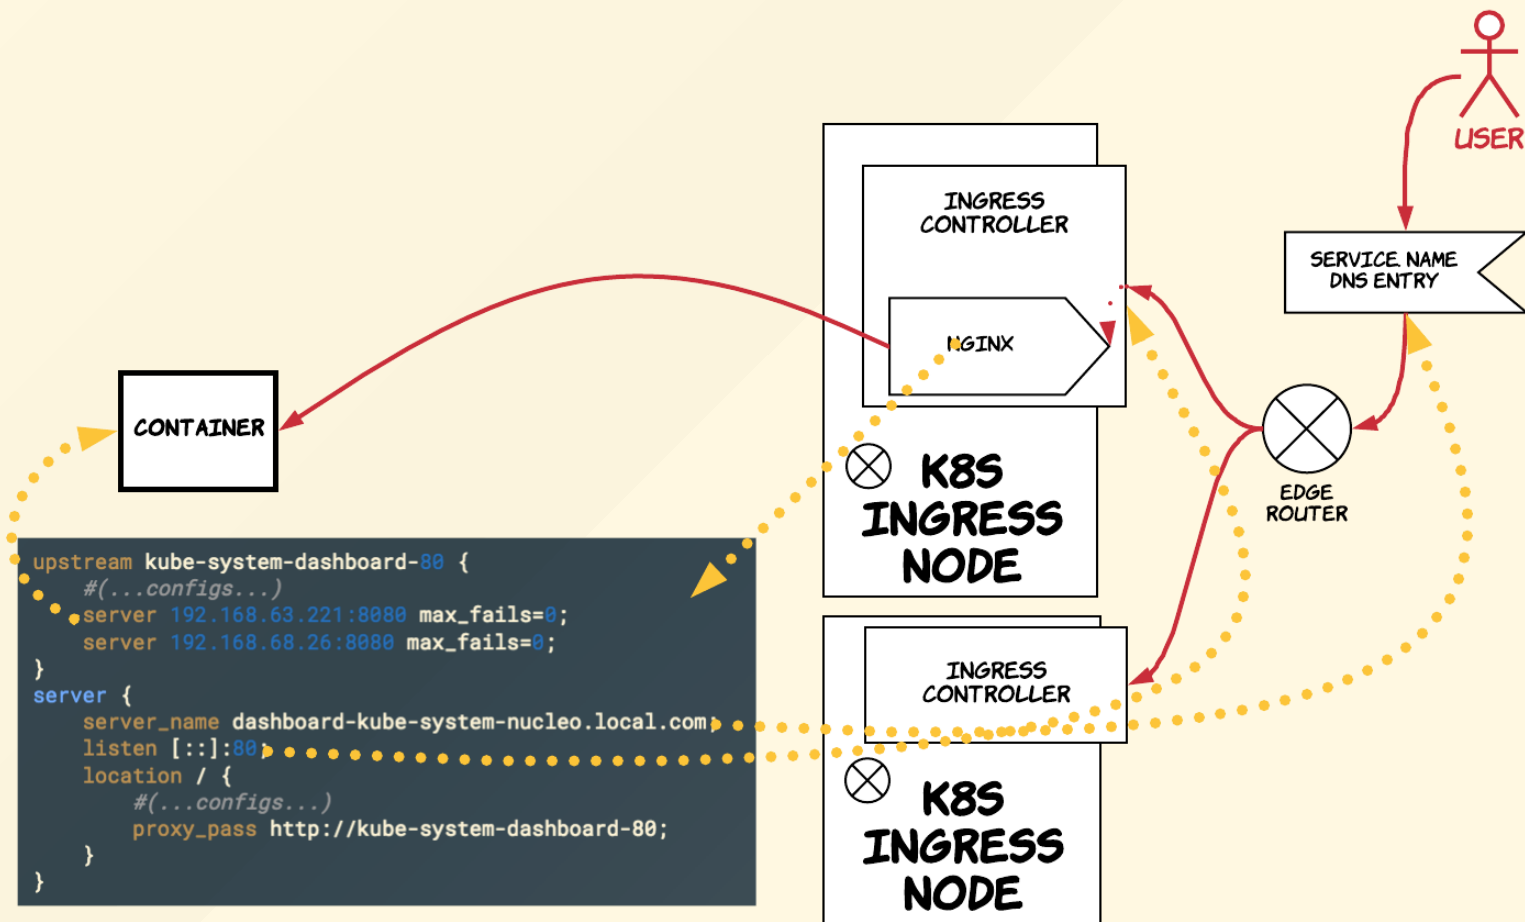
```
$ kubectl get pods -n al-lee -l app=nucleo,release=al-lee-fernet-decryptor-0162733
NAME                                READY   STATUS    RESTARTS   AGE
fernet-decryptor-0162733-5b694965f6-zdSz8   1/1     Running   0           3d
```

**BUT, USE POD IP NOT CLUSTERIP WHICH USE NAT
THIS IS ONLY FOR SERVICE DISCOVERY**

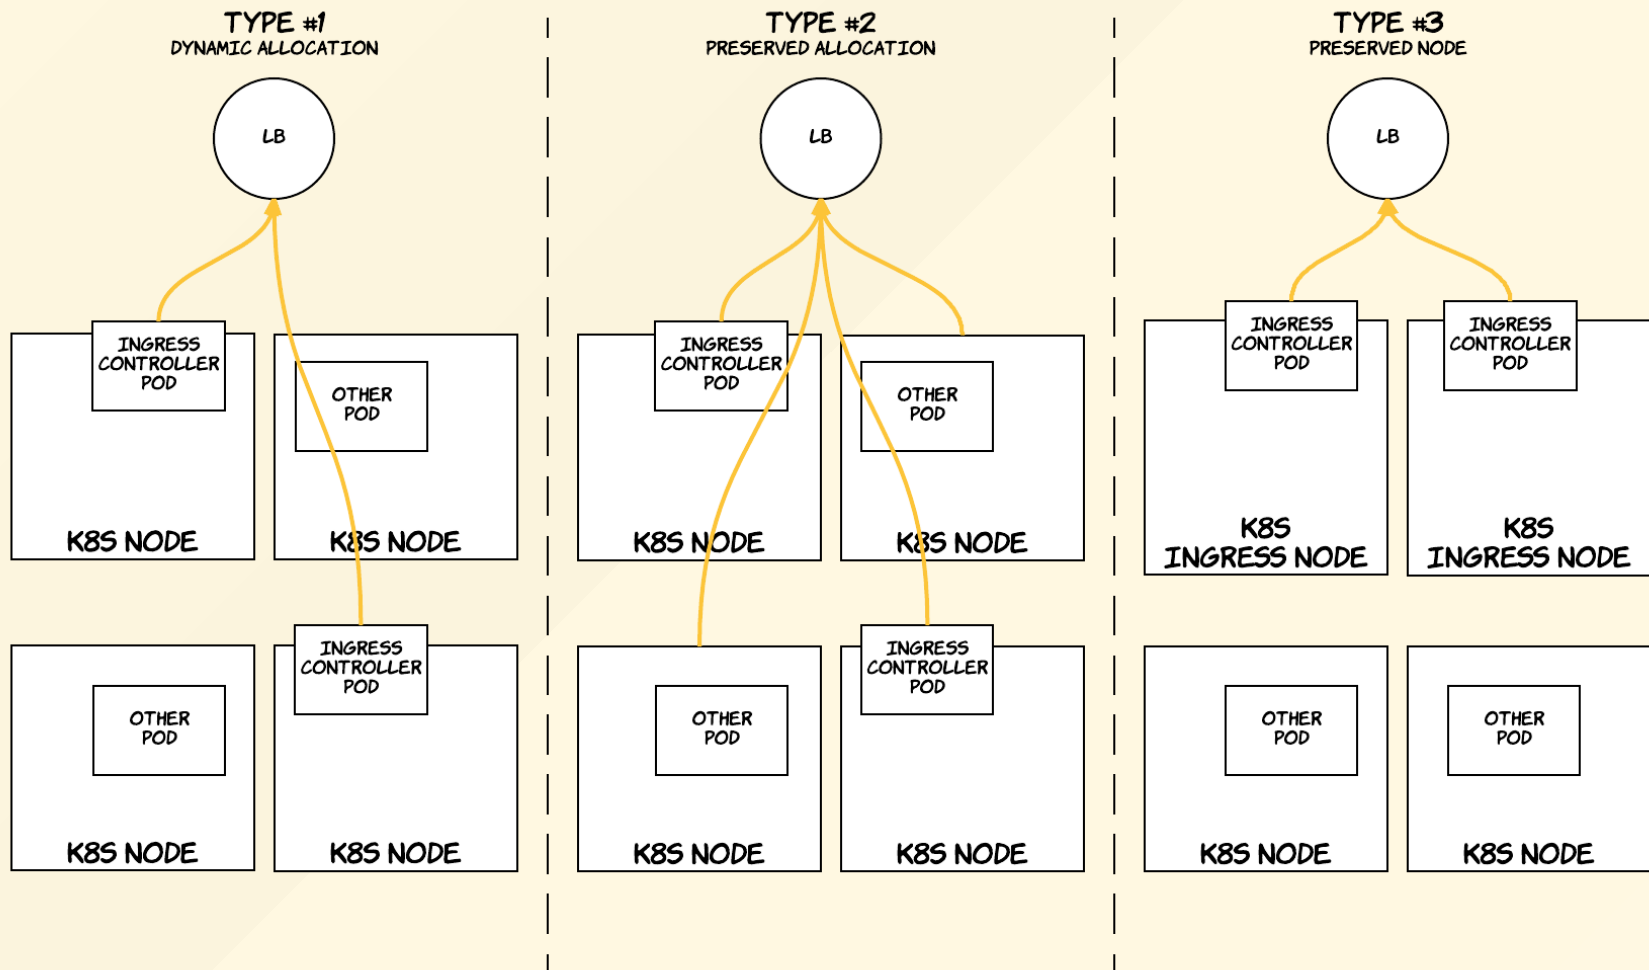
Ingress with controller



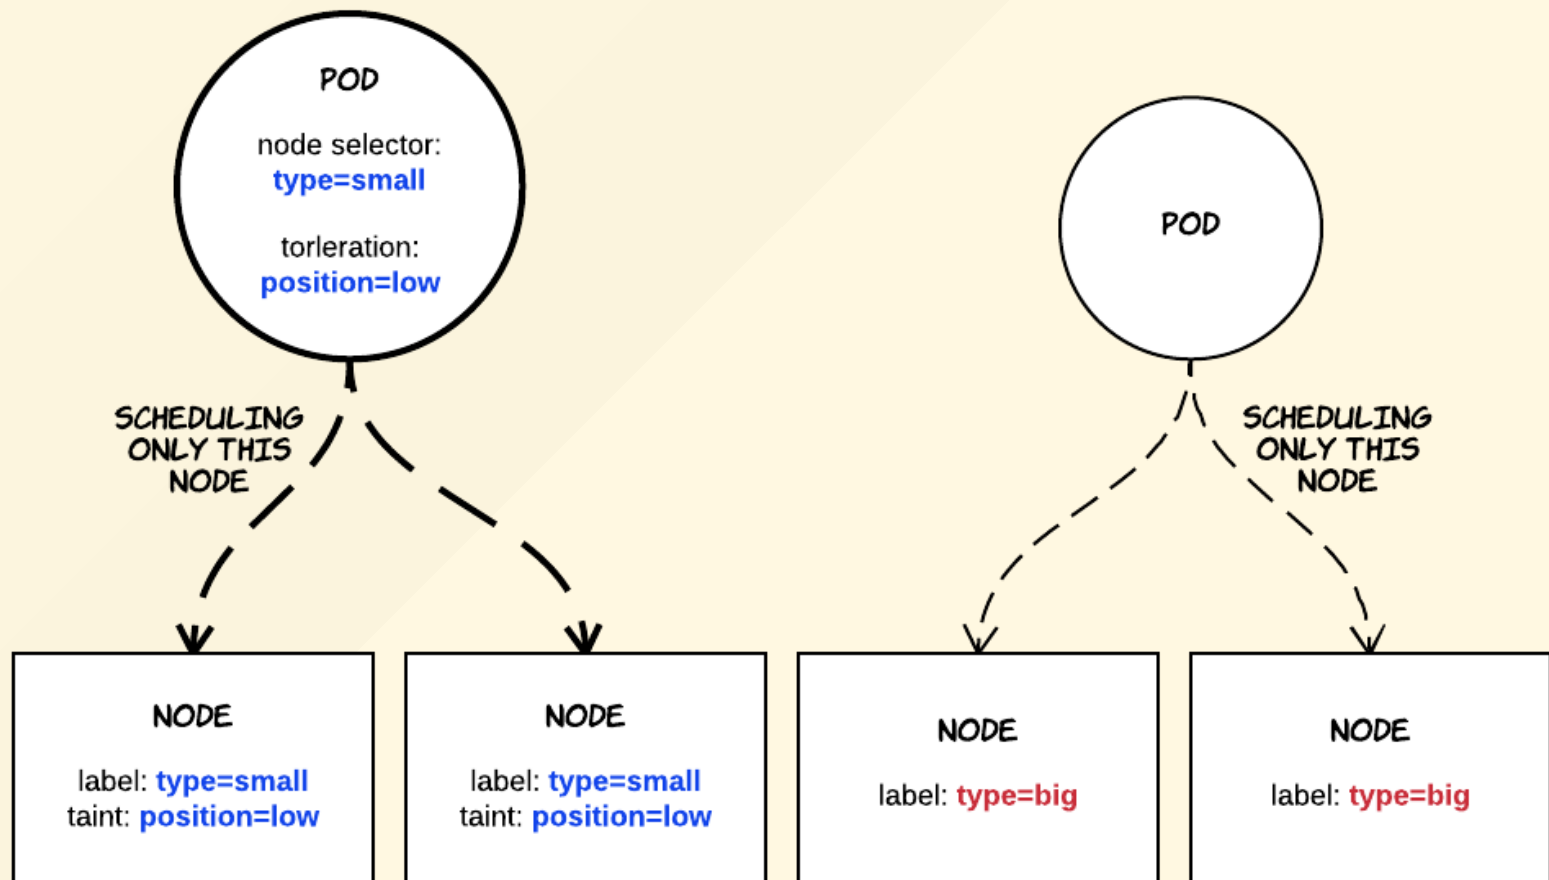
Ingress



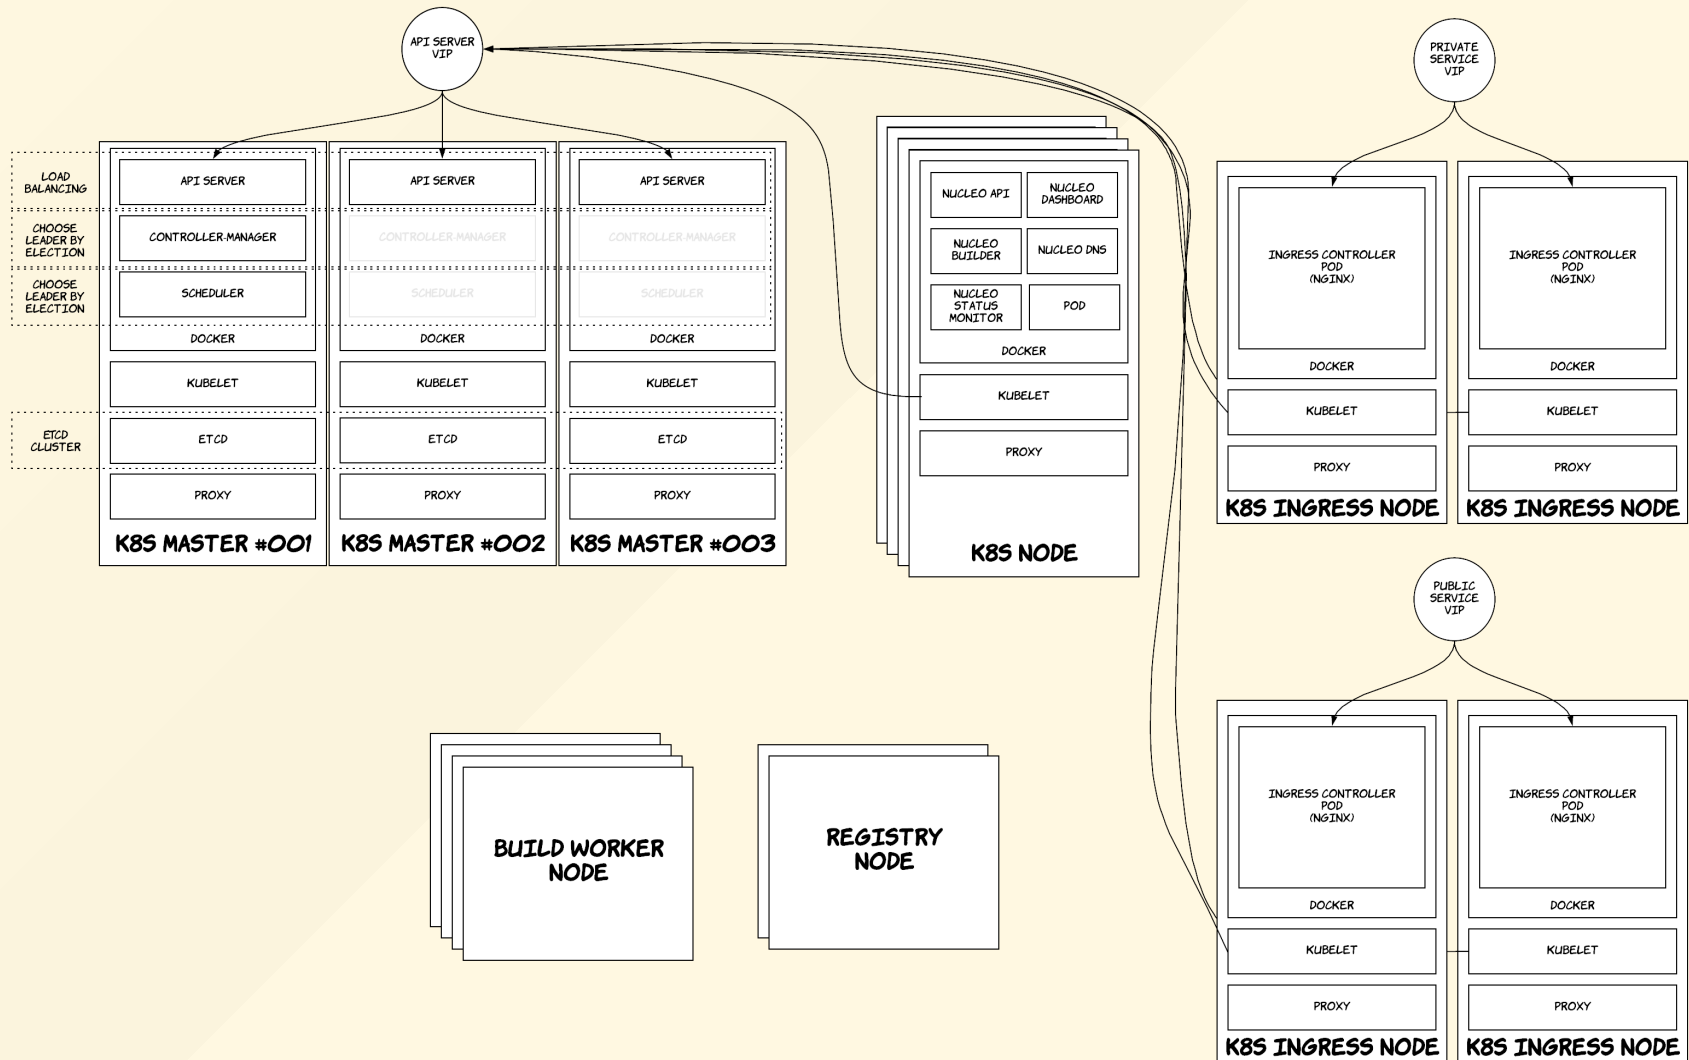
Ingress node deploy types



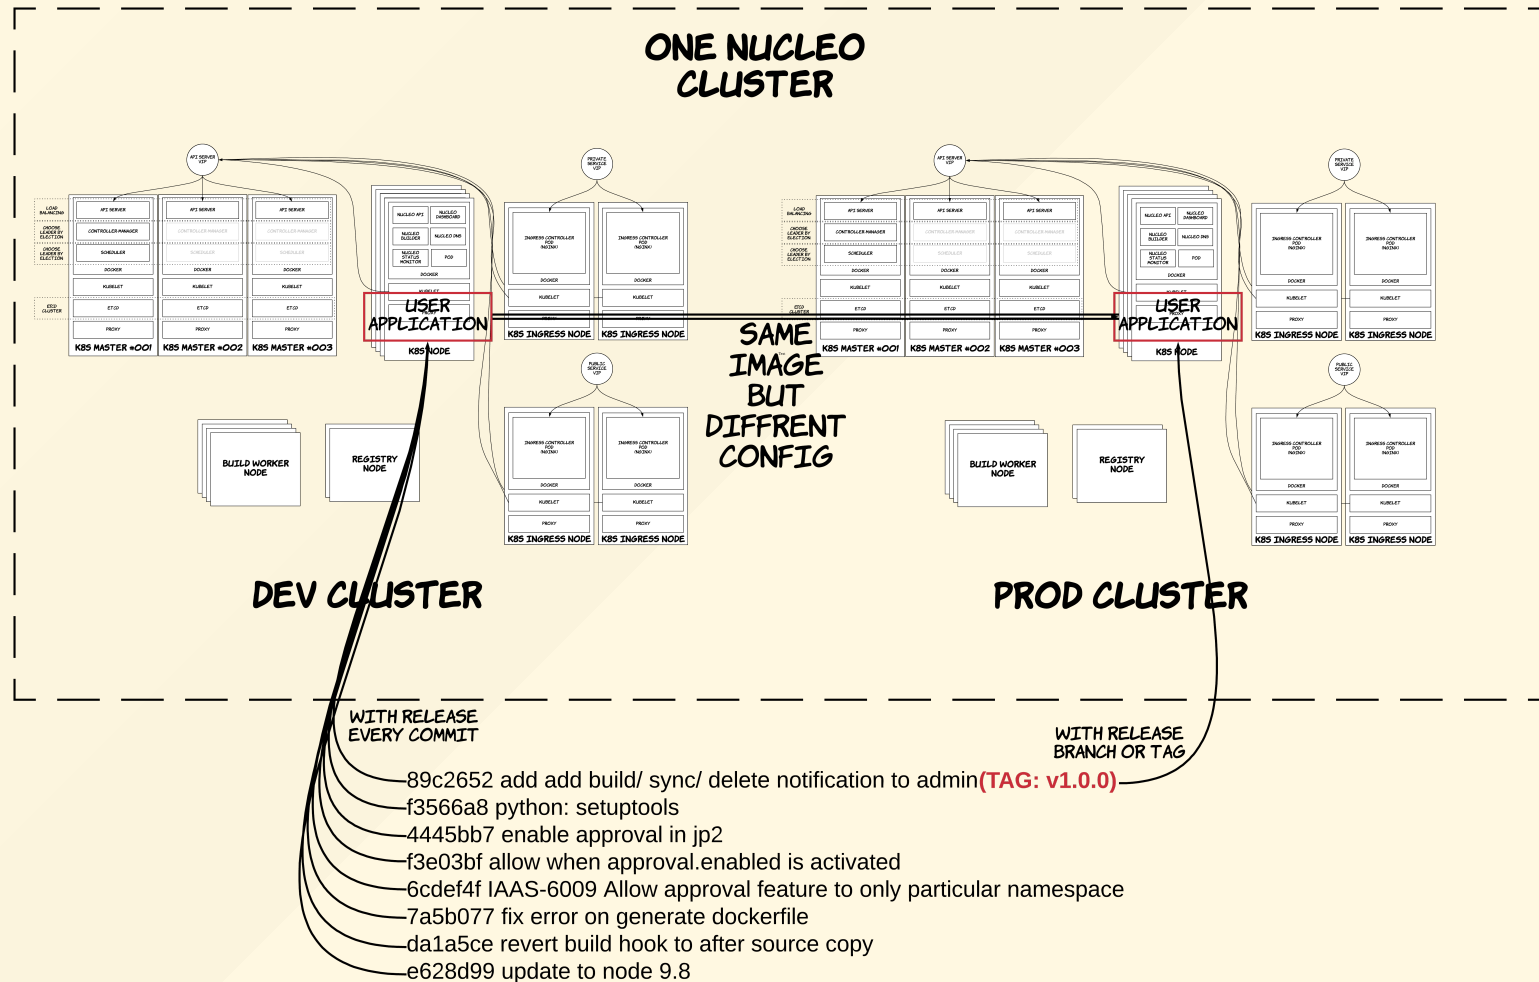
Dedicated node pattern for ingress



Update k8s & nucleo architecture



How to support continuous delivery



- currently not use kubernetes federation

winner winner chicken dinner

Notify kubernetes-dev of lifting code freeze	Lead			20			
v1.10.0	Branch Manager		3/26		week 13	1.10-blocking	
v1.11.0-alpha.1	Branch Manager			27			
Release retrospective	Community			29			
1.11 Release Cycle Begins	Next Lead				2		

Thank you